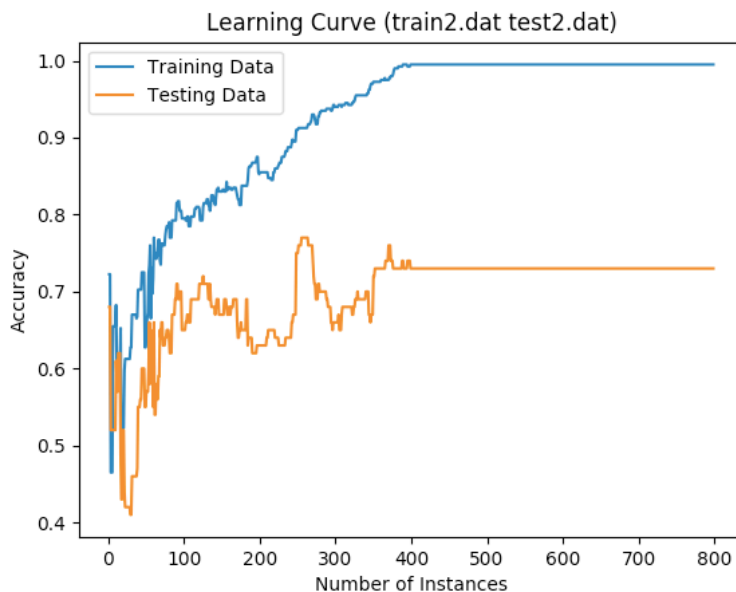




This graph shows Number of Instances vs. Accuracy for the train.dat and test.dat files.

The shape of these curves is interesting. We see that testing data pretty much always tests worse than training data.

Notably, the curves simply level off in accuracy rather than degrade as we might expect. We expected training data to continue to increase in performance and for testing data to reduce after peaking.



This graph is similar, except it is the plot using train2.dat and test2.dat

These curves are somewhat closer to what we expect. The testing data performance does have a small peak around $X = 275$ and then performance somewhat degrades. There is also much worse testing performance between the testing and training data.